

**LECTURE NOTES IN CIS300**

**YUZHE (RICHARD) TANG**

**SPRING, 2018**

# **SECTION 2: C/C++ PROGRAMMING**

# REFERENCES

- "Unix Programming Tools", [[link](#)]
- Computer Systems: A Programmer's Perspective, Randal E. Bryant and David R. O'Hallaron, Chapter 1, [[online pdf](#)]

# HELLOWORLD C

```
#include <stdio.h> //include header file for preprocessing
int y = 3; //initializing a global variable
//extern int y; //declaring global variable
int main() //defining main function
{
    int x = 0; //local var. (def. & init.), literal,
    printf("helloworld: y = %d\n",y); //function (invocation)
    return 0;
}
```

- `printf` is a function
  - the first argument is *format string*
- header files: `stdio.h`

# LIFE CYCLE OF A VARIABLE/FUNCTION

|            | variable                   | function                          |
|------------|----------------------------|-----------------------------------|
| declare    | <code>extern int x;</code> | <code>void foo( );</code>         |
| define     | <code>int x;</code>        | <code>void foo( ){ }</code>       |
| initialize | <code>int x=6;</code>      |                                   |
| reference  | <code>y=x;x=1;</code>      | <code>foo( );</code> (invocation) |
| destroy    |                            |                                   |

# COMPILATION & EXECUTION: BASICS

- GCC: GNU Compilation Collection
- In your terminal, run the following commands

```
gcc hello.c  
./a.out
```

# EXERCISES

- Write a C program that prints out your name. Compile and execute it in Ubuntu. Submit the C program to BB.
- Write a C program that computes the sum of 1,2,3,...,956. Compile and execute the program in Ubuntu. Submit the C program to BB.

**GCC**




# COMPILATION (1)

- Two steps of compilation:
  - *compiling*: text `.c` file to relocatable `.o` (object) file
  - *linking*: multiple relocatable `.o` files to one executable `.o` file
    - *symbol*: reference to link construct (declaration) in one `.o` file to construct (definition) in another `.o` file

# COMPILATION (2)

```
gcc hello.c -o a.out  
gcc -S hello.c -o hello.s #compiler  
gcc -c hello.s -o hello.o #assembler  
gcc hello.o -o a.out #linker
```

- gcc is a compilation system
  1. **preprocessor**: from source file to source
  2. **compiler**: from source to assembly file
    - *assembly file*
  3. **assembler**: from assembly file to relocatable object file
  4. **linker**: from multiple objects to an executable object



Linker

# COMPILING MULTIPLE C PROGRAMS

In file1.c:

```
#include <stdio.h>
extern void foo();
int main(){
    printf("main();\n");
    foo();
}
```

In file2.c:

```
#include <stdio.h>
void foo(){
    printf("foo();\n");
}
```

# COMPILING MULTIPLE C PROGRAMS (2)

```
gcc file1.c file2.c  
# try this?  
gcc file1.c  
gcc file2.c
```

# COMPILING MULTIPLE C PROGRAMS (3)

```
gcc -c file1.c # compiler & assembler  
gcc -c file2.c # compiler & assembler  
gcc file1.o file2.o # linker
```

Or

```
gcc -S file1.c # compiler  
gcc -c file1.s # assembler  
gcc -S file2.c # compiler  
gcc -c file2.s # assembler  
gcc file1.o file2.o # linker
```

# LINK LIBRARY FILES

```
gcc -S file1.c # compiler  
gcc -c file1.s # assembler  
gcc file1.o file2.o # linker
```

```
mv file2.o ../libfile2.a  
gcc file1.o ../libfile2.a # linker  
gcc file1.o -L.. file2.o # linker  
gcc file1.c -L.. file2.o # linker
```

- Gcc flag: `-Ldir -lmylib` for library to link

# INCLUDE HEADER FILE

In header1.h:

```
extern void foo();
```

In file1.c:

```
#include <stdio.h>
#include "header1.h"
int main(){
    printf("main();\n");
    foo();
}
```

```
gcc file1.c file2.c
```



# INCLUDE HEADER FILE (2)

Header file in another directory

```
mv header1.h ..  
#will the following work?  
gcc file11.c file2.c  
gcc -I .. file11.c file2.c
```

- Gcc flag for searching header file: `-I path`
  - `path` is where the header file is

# GCC FLAGS (SUMMARY)

- `-c` for specifying using gcc as a compiler
- `-o` for specifying the name of output file
- `-Ldir -lmylib` for linking a library
  - search library file `dir/libmylib.a` for solving symbols (functions, global variables) when linking
- `-I` for `#include`
  - header file (storing declarations)
- `-Wall, -W` for warning
- `-g` for debug (later): `gcc -g file1.c file2.c`
- ref [[link](#)]

# EXERCISE

- Write two C files:
  - `filea.c` defines functions `main( )` and `bar( )`
  - `fileb.c` defines function `foo( )`
  - function `main( )` calls `foo( )`
  - function `foo( )` calls `bar( )`
  - Compile your program.
  - Submit the program and commands to BB.

# MAKE AND MAKEFILE

# DOWNLOAD COURSE REPO.

To download course repository, type the following commands

```
sudo apt-get update
sudo apt-get upgrade
sudo apt-get install git
git clone https://github.com/SUCourses/cis300-18spring.git
```

# MAKEFILE: DEPENDENCY RULES

- `make` is a tool for project management in shell
- `Makefile` is the configuration file that tells `make` what to do
- A `Makefile` consists of a series of dependency rules
- Each dependency rule expresses IFTTT logic (if-this-then-that)

```
target: files/objects  
(tab)commands
```

There is a **tab** before the commands

# HELLOWORLD MAKEFILE

In Makefile (related files are in dir. demos/mar7)

```
all:
    gcc file1.c file2.c
```

To run it, type following command in a terminal

```
make
```

(Try change `file.c`, and make it again).

# MAKEFILE OF MULTIPLE RULES

```
c:
    gcc file1.c file2.c

exec: c
    ./a.out

clean:
    rm *.o *.out
```

Note there are empty lines btwn. two rules.



# USE MAKEFILE TO LINK (1)

Recall how to run compiler, assembler and linker

```
gcc -c file1.c # compiler & assembler  
gcc -c file2.c # compiler & assembler  
gcc file1.o file2.o # linker
```

# USE MAKEFILE TO LINK (2)

A Makefile can easily manage different "targets" - It does compiling, assembling and linking separately

```
link: file1.o file2.o
    gcc file1.o file2.o

file1.o: file1.c
    gcc -c file1.c

file2.o: file2.c
    gcc -c file2.c
```

```
make
make
```

# USE MAKEFILE TO LINK (3)

Use a default rule to compile individual C file

```
link: file1.o file2.o
    @gcc file1.o file2.o
```

```
make
make
```

- @ used to hide the command in printout.

# MAKEFILE: USING VARIABLES

```
SRCS = file1.c file2.c
OBJS = $(SRCS:.c=.o)
CFLAGS = -g -Iheaders
#LDFLAGS = -L. -lxxx

link: $(OBJS)
      $(CC) $(LDFLAGS) $(OBJS)
```

# MAKEFILE: USING VARIABLES (2)

- A Makefile variable is a text string
- There're standard variables
  - `CC` is equal to `gcc` (the compiler)
  - `OBJS = $(SRCS:.c=.o):`
    - This incantation says that the object files have the same name as the `.c` files, but with `.o` extension
  - `LDFLAGS` is the linker flag of library search path (`-L`)
  - `CFLAGS` is the default compiling flags

# EXERCISE

1. Write a `Makefile` such that `make` always clean `.o` files, recompiles all `.c` files and executes the new `.o` file.
2. Write a `Makefile` such that `make link` will compile a `file.c` file against a library file `libxxx.a`

**GDB**

# REFERENCES

- "Reviewing gcc, make, gdb, and Linux Editors", [[pdf](#)]
- "Unix Programming Tools", [[link](#)]



# A BUGGY C PROGRAM

```
#include<stdio.h> //printf
int array_stack[] = {0,1,2};
int main(){
    int sum; // local variable
    for(int i=0; i<=3; i++){
        sum += array_stack[i];
    }
    printf("sum = %d\n", sum);
    return 0;
}
```

# USE GDB TO FIND BUG

- Installing gdb
  - on MacOS: [[youtu.be/Vj33vsrDkE80](https://youtu.be/Vj33vsrDkE80)]
  - on Ubuntu: `sudo apt-get install gdb`
- Compile: `gcc -g`
- Run gdb: `gdb a.out`

# GDB COMMAND: CONTROL EXECUTION

- CPU executes a C program, statement by statement
- Breakpoint is a GDB mechanism to control where CPU should pause execution
  - `break/b` is a GDB command to set breakpoint.
  - `break/b file:n|fn|file:fn`: breakpoint can be file:line number (n), function name or file:function name.
  - `disable/enable/delete` are GDB commands to disable/enable/delete a breakpoint
  - `disable/enable/delete i`: i is the index of breakpoint

# GDB COMMAND: CONTROL EXECUTION

## (2)

- Stepping is a GDB mechanism to control the CPU execution step by step.
  - `run / r`: start to run the program until next breakpoint or the end of program.
  - `next / n`: run just the next statement (step over a function call)
  - `continue / c`: continue the execution until the next breakpoint

# GDB COMMAND: EXAMINE RUNTIME

- Examine runtime data
  - `print v/p` `v`: print variable `v`
- Examine code (with `gcc -g`)
  - `list/l`
- Examine execution environment: e.g. stack (later)

# GDB COMMANDS

| functionality | commands   |
|---------------|--|
| breakpoints   | b,disable/enable/delete breakpoi                           |
| stepping      | r,s,n,c,finish,return                                      |
| examine_data  | p/i v,display/undisplay,watch,set                          |
| examine_code  | list   |
| examine_stack | bt,where,info,up/down,frame                                |
| misc.         | editmode vi,b fn if expression,h<br>disassembler,shell cmd |

# DEMO

- Debug the following program using gdb

```
#include<stdio.h> //printf
int array_stack[] = {0,1,2};
int main(){
    int sum; // local variable
    for(int i=0; i<=3; i++){
        sum += array_stack[i];
    }
    printf("sum = %d\n", sum);
    return 0;
}
```

# EXERCISE

- Exercise: Debug the following program using gdb, upload the correct program to BB.

```
#include<stdio.h>
int main() {
    int x = 5;
    int y = 3;
    int z = x - y;
    int a = x * y;
    int b = a - 7*z;
    b--;
    int c = z + y;
    int d = c / b;
    int e = a + 12;
    int f = e - b;
    printf("%d\n",f);
}
```



# POINTER IN C

# REFERENCES

- Pointer Basics: [<http://cslibrary.stanford.edu/106/>]
- Point fun with Binky: [<http://cslibrary.stanford.edu/104/>]

# C POINTER

- A C pointer is a C variable that stores the reference to something.
  - "something", called pointee, is usually another variable.
- In the figure below, a pointer variable named `x` stores a reference to a "pointee" variable of value 42.



pointer pointee

# POINTER OPERATIONS

- Definition/initialization: `int *p1 = p2;`
- Assignment: `p1 = p2;`
- Dereference: `*p = 3`
- Get reference of: `&a`
  - get the *address* (memory location) of variable a

```
#include<stdio.h>
int main(){
    int a = 10;
    int * p = & a;
    int b = *p;
    printf("a=%d,b=%d,*p=%d,p=%p\n",a,b,*p,p);
}
```

# BINKY'S CODE (1)

```
void main() {  
    int*    x;  // Allocate the pointers x and y  
    int*    y;  // (but not the pointees)  
}
```



Allocate pointer

# BINKY'S CODE (2)

```
void main() {  
    int*    x;  // Allocate the pointers x and y  
    int*    y;  // (but not the pointees)  
    x = malloc(sizeof(int));    // Allocate an int pointee,  
                                // and set x to point to it  
}
```



Allocate pointee

# BINKY'S CODE (3)

```
void main() {  
    int*    x;  // Allocate the pointers x and y  
    int*    y;  // (but not the pointees)  
    x = malloc(sizeof(int));    // Allocate an int pointee,  
                                // and set x to point to it  
    *x = 42;    // Dereference x to store 42 in its pointee  
}
```



Dereference pointer

# BINKY'S CODE (4)

```
void main() {  
    int*    x;  // Allocate the pointers x and y  
    int*    y;  // (but not the pointees)  
    x = malloc(sizeof(int));    // Allocate an int pointee,  
                                // and set x to point to it  
    *x = 42;    // Dereference x to store 42 in its pointee  
    *y = 13;    // CRASH -- y does not have a pointee yet  
}
```




Dereference failure



# BINKY'S CODE (5)


```
void main() {  
    int*    x;  // Allocate the pointers x and y  
    int*    y;  // (but not the pointees)  
    x = malloc(sizeof(int));    // Allocate an int pointee,  
                                // and set x to point to it  
    *x = 42;    // Dereference x to store 42 in its pointee  
    *y = 13;    // CRASH -- y does not have a pointee yet  
    y = x;      // Pointer assignment sets y to point to x's pointee  
}
```



Pointer assignment

# BINKY'S CODE (6)

```
void main() {  
    int*    x;  // Allocate the pointers x and y  
    int*    y;  // (but not the pointees)  
    x = malloc(sizeof(int));    // Allocate an int pointee,  
                                // and set x to point to it  
    *x = 42;    // Dereference x to store 42 in its pointee  
    *y = 13;    // CRASH -- y does not have a pointee yet  
    y = x;      // Pointer assignment sets y to point to x's pointee  
    *y = 13;    // Dereference y to store 13 in its (shared) pointee  
}
```



Deference pointer

# LIFE CYCLE OF A C POINTER

|               | pointer                      | variable                  | function          |
|---------------|------------------------------|---------------------------|-------------------|
| declare       | <code>extern int * p</code>  | <code>extern int x</code> | <code>void</code> |
| define        | <code>int *p;</code>         | <code>int x</code>        | <code>void</code> |
| initialize    | <code>int *p=&amp;a;</code>  | <code>int x=6</code>      |                   |
|               | <code>int*q=malloc(7)</code> |                           |                   |
| (de)reference | <code>*p=x; x=*p</code>      | <code>y=x</code>          | <code>for</code>  |
| destroy       | <code>delete p</code>        |                           |                   |

# EXERCISE

- Do the following to complete the code snippet at the bottom. Then compile and execute your program. Submit the completed program to BB.
  1. define two pointers `p1` and `p2`, both pointing to variable `x`.
  2. Use `p1` to update `x`'s value to 5.
  3. Then use `p2` to read the value of variable `x` and `printf` it on terminal.

```
#include<stdio.h>
int main(){
    int x = 4;
    // To complete the program below:

}
```

# **C POINTER AND DATA TYPES**

# DATA TYPE

- C is a typed language
- Data type in C determines:
  - How much space to allocate for storing a variable in memory
  - How to interpret bit-string stored in the memory
  - How to carry out the arithmetic on primitive types

# PRIMITIVE TYPES

- types: signed, unsigned, long long, float, char

| type                  | signed | unsigned | short | long<br>long | float | char |
|-----------------------|--------|----------|-------|--------------|-------|------|
| <code>sizeof()</code> | 4      | 4        | 2     | 8            | 4     | 1    |

- unsigned: a 32-bit unsigned integer, value from 0 to  $2^{32} - 1$ .
- signed: a 32-bit signed integer, value from  $-2^{31}$  to  $2^{31} - 1$ .
  - first bit determines whether negative
- Typecasting: convert the type of a variable.
  - `int x = 1; double f = (double) x;`

# DEMO 1: TYPE INTERPRETATION

```
#include<stdio.h>
int main(){
    unsigned int u = 2147483649;
    int v = (int) u;
    printf("unsigned vs signed: %ud,%d\n",u,v);

    int i=1;
    float f = (float) i;
    printf("float vs int: %f,%d\n",f/3,i/3);
}
```



# DEMO 2: DATA TYPE SIZE

```
#include<stdio.h>
int main(){
    signed int a;
    unsigned int b;
    short c;
    long long d;
    float e;
    char f;
    printf("signed int: %lu\n", sizeof(a));
    printf("unsigned int: %lu\n", sizeof(b));
    printf("short: %lu\n", sizeof(c));
    printf("long long: %lu\n", sizeof(d));
    printf("float: %lu\n", sizeof(e));
    printf("char: %lu\n", sizeof(f));
    return 0;
}
```

# POINTER AND ARRAY

- An array in C stores a list of elements in adjacent memory locations.
- Use pointer to access array element
  - Pointer type: `char *`, `int *`
  - Pointer arithmetic:
    - `int * p = array; p += 1;`
    - `int pp = array; pp += sizeof(int);`

```
#include<stdio.h>
int main(){
    int a[] = {2,1,0};
    int *b = a; // b points to the first element in a
    unsigned long c = (unsigned long)a;//long
    for (int i=0; i<3; i++){
        printf("%d,%d,%d,%d,%d\n",a[i],*(b+i),*(a+i),b[i],*((int *)(&c+i)))
    }
}
```

# FUNCTION POINTER

- Two classes of pointers
- Data pointer: pointer to variable, array
- Code pointer: function pointer

```
#include <stdio.h>  /* for printf */  
// https://en.wikipedia.org/wiki/Function\_pointer  
double cm_to_inches(double cm) {  
    return cm / 2.54;  
}  
int main(void) {  
    double (*func1)(double) = cm_to_inches;  
    printf("%f\n", func1(15.0));  
    return 0;  
}
```

# EXERCISE

1. Write a C program that defines function `void foo(void)` and `int bar(long x)`. Call these two functions through function pointers. Upload your program to BB.
2. Complete the following program that scans the array using index `long_index`. Upload your program to BB.

```
#include<stdio.h>
int main(){
    int a[] = {7,9,6};
    unsigned long long_index = (unsigned long)a;
    for(int i=0; i<3; i++){
        printf("%d,",*(int*)(long_index));
        long_index += XXX; // fill out XXX
    }
}
```

**FILE I/O**

# REFERENCES

- "Advance Programming in the Unix Environment" (APUE), Chapter 3.1-3.8 [[link](#)]

# INTRODUCTION

- Five functions: `open`, `read`, `write`, `lseek`, `close`
- They are unbuffered IO in the sense that each call (`read`) invokes a syscall.
  - Unbuffered IO functions are not ISO C, but part of POSIX.1.
- Atomic functions over shared resources.

# FILE DESCRIPTORS

- All open files are referred to by file descriptors.
- A file descriptor is a non-negative integer.
- FD is returned by `open` or `creat`, and is used as argument to `read` or `write`.
- 0 is FD for `stdin`, 1 is the FD for `stdout`, 2 is FD of `stderr`.



# OPEN

## Open a file

```
#include <fcntl.h>
int open(const char * pathname, int oflag, ...);
//returns: file descriptor if OK, -1 on error
```

- `oflag` takes one of three mandatory values and OR with optional values.
  - mandatory: `O_RDONLY`, `O_WRONLY`, `O_RDWR`
  - optional:
    - `O_CREAT`: Create a file if it doesn't exist
    - `O_TRUNC`: Truncate a file to zero if it exists and if it is opened for write-only or read-write
- the file descriptor returned is the lowest-numbered unused descriptor.

# CREAT

Create a file

```
#include <fcntl.h>
int creat(const char * pathname, mode_t mode);
//returns: file descriptor opened for write-only if OK, -1 on error
//equiv. to
open(pathname, O_WRONLY | O_CREAT | O_TRUNC, mode);
```

# CLOSE

Close a file

```
#include <unistd.h>  
int close(int fd);  
//return: 0 if OK, -1 on error
```

# LSEEK

- Every open file has a "current file offset"
- Read and write starts at the offset and cause it to increment by the number of bytes read/written.

```
#include <unistd.h>
off_t lseek(int fd, off_t offset, int whence);
//Returns: new file offset if OK, 1 on error
```

- whence:
  - SEEK\_SET: set offset to be offset plus the beginning of the file.
  - SEEK\_CUR: set offset to offset plus the current value.
  - SEEK\_END: set offset to be file size plus offset

# READ AND WRITE

```
#include <unistd.h>
ssize_t read(int fd, void *buf, size_t nbytes);
//Returns: number of bytes read, 0 if end of file, 1 on error
```

- It requests to read `nbytes` bytes from `fd` and stores them in `buf`.
- If the read is successful, it returns the actual number of bytes read.
- If the end of a file is reached, it returns 0

```
#include <unistd.h>
ssize_t write(int fd, const void *buf, size_t nbytes);
//Returns: number of bytes written if OK, 1 on error
```

- It requests to write `nbytes` bytes to `fd` from `buf`.

# SEEKABLE FILES

```
#include<unistd.h>
#include<stdio.h>
#include<stdlib.h>

int main(void){
    if (lseek(STDIN_FILENO, 0, SEEK_CUR) == -1)
        printf("cannot seek\n");
    else
        printf("seek OK\n");
    exit(0);
}
```

```
> ./a.out
> ./a.out < file #redirection is seekable
> cat file | ./a.out #pipe file is not seekable
```

# CREAT FILE WITH A HOLE

- header.h

```
#include<stdio.h>
#include<unistd.h> //lseek, STDIN_FILENO
#include<stdlib.h>
#include <fcntl.h>
#define FILE_MODE (S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH)
void err_sys(const char* x) {
    perror(x);
    exit(1);
}
```

```

#include "header.h"
char buf1[] = "abcdefghij";
char buf2[] = "ABCDEFGHIJ";
int main(void){
    int fd;
    if ((fd = creat("file.hole", FILE_MODE)) < 0)
        err_sys("creat error");
    if (write(fd, buf1, 10) != 10)
        err_sys("buf1 write error"); /* offset now = 10 */
    //comment out the following two lines, get a file file.nohole
    if (lseek(fd, 16384, SEEK_SET) == -1)
        err_sys("lseek error"); /* offset now = 16384 */
    if (write(fd, buf2, 10) != 10)
        err_sys("buf2 write error"); /* offset now = 16394 */
    exit(0);
}

```

```

> cat file.hole
> cat file.nohole
> ls -ls file.hole file.nohole
> od -c file.hole
> od -c file.nohole

```



# EXERCISE

- Write a C program that does the same thing to the following shell script. Upload your code to BB.

```
touch file1.txt  
echo "Alice" >> file1.txt  
cat file1.txt
```

Hint: to print a char array buf, `printf( "%s\n", buf );`

## **FILE I/O (2)**

# READ AND WRITE (CONT'ED)

- Demo1: Copy text from stdin to stdout
  - header.h

```
#include<stdio.h>
#include<unistd.h> //lseek, STDIN_FILENO
#include<stdlib.h>
#include <fcntl.h>
#define FILE_MODE (S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH)
void err_sys(const char* x) {
    perror(x);
    exit(1);
}
```


```
#include "header.h"
#define BUFFSIZE 4096
int main(void){
    int n;
    char buf[BUFFSIZE];

    while ((n = read(STDIN_FILENO, buf, BUFFSIZE)) > 0)
        if (write(STDOUT_FILENO, buf, n) != n)
            err_sys("write error");

    if (n < 0)
        err_sys("read error");
    exit(0);
}
```

# FILE SHARING & ATOMIC OPERATIONS

- An open file can be shared among multiple processes
- There are operations that are safe in multi-processing environment (atomic operations)
  - File append
  - `pread/pwrite` function
  - `sync, fsync`



File kernel representation and sharing

# FILE APPEND

- Two ways to append text to a file: `lseek/write` and `O_APPEND`.
  - `lseek/write` is not safe in shared files.
  - `O_APPEND` is safe.

```
if (lseek(fd, 0L, SEEK_END) < 0) /* position to EOF */  
    err_sys("lseek error");  
if (write(fd, buf, 100) != 100) /* and write */  
    err_sys("write error");
```

```
open(pathname, O_WRONLY | ... | O_APPEND, mode);
```

# PREAD/PWRITE FUNCTION

- `pread` and `pwrite` functions: do seek and perform I/O atomically.
  - Calling `pread/pwrite` is equivalent to calling `lseek` followed by a call to `read/write`. There is no way to interrupt `lseek` and `read/write`.

```
#include <unistd.h>
ssize_t pread(int filedes, void *buf, size_t nbytes, off_t offset);
// Returns: number of bytes read, 0 if end of file, 1 on error

ssize_t pwrite(int filedes, const void *buf, size_t nbytes, off_t offset);
//Returns: number of bytes written if OK, 1 on error
```



# **SYNC, FSYNC**

- Traditional UNIX system has a page cache in the kernel that queues data writes before the buffer overflows and it flushes data to disk.
- `sync` / `fsync` ensures consistency of the file system on disk with the contents of the buffer cache.
- The `sync` function simply queues all the modified block buffers for writing and returns; it does not wait for the disk writes to take place.
- The function `fsync` waits for the disk writes to complete before returning. The intended use of `fsync` is database applications that need be sure that the modified blocks have been written to the disk.

# FCNTL

- The `fcntl` function reads/changes the properties of an open file.

```
#include <fcntl.h>
int fcntl(int filedes, int cmd, ... /* int arg */ );
//Returns: depends on cmd if OK (see following), 1 on error
```

# DEMO2: PRINT FLAGS FOR SPECIFIED DESCRIPTOR

- see `demos/apr16/demo2.c`

```
#include "header.h"
int main(int argc, char *argv[]) {
    int val;
    if (argc != 2)
        err_quit("usage: a.out <descriptor#>");
    if ((val = fcntl(atoi(argv[1]), F_GETFL, 0)) < 0)
        err_sys("fcntl error for fd %d", atoi(argv[1]));
    switch (val & O_ACCMODE) {
    case O_RDONLY:
        printf("read only");
        break;
    case O_WRONLY:
        printf("write only");
        break;
    case O_RDWR:
        printf("read write");
```

```
./a.out 0 < /dev/tty
./a.out 1 > temp.foo; cat temp.foo
./a.out 2 2>> temp.foo
```

```
./a.out 5 5<> temp.foo
```

5<>temp.foo opens file temp.foo for read/write on file descriptor 5

# STAT

- `stat` shell command
- `stat` returns information about a named file.
- `fstat` returns information about the open file by descriptor.

```
#include <sys/stat.h>

int stat(const char *restrict pathname, struct stat *restrict buf);
int fstat(int fd, struct stat *buf);
//All three return: 0 if OK, 1 on error
```

# FILE TYPES

- Regular file: The most common file type, which contains data of some form.
- Directory file: file that contains the names of other files and pointers to information on these files.
- Block special file: file providing buffered I/O access in fixed-size units to devices such as disk drives.
- Character special file: file providing unbuffered I/O access in variable-sized units to devices.
- FIFO: file used for communication between processes, such as named pipe.
- Socket: file used for network communication between processes.
- Symbolic link: file that points to another file.

# DEMO3: PRINT FILE TYPES

```
int main(int argc, char *argv[]) {
    int i;
    struct stat buf;
    char *ptr;

    for (i = 1; i < argc; i++) {
        printf("%s: ", argv[i]);
        if(lstat(argv[i], &buf) < 0) {
            err_ret("lstat error");
            continue;
        }
        if (S_ISREG(buf.st_mode))
            ptr = "regular";
        else if (S_ISDIR(buf.st_mode))
            ptr = "directory";
        else if (S_ISCHR(buf.st_mode))
```

```
>./a.out /etc/passwd
/etc/passwd: regular
>./a.out .
.: directory
>./a.out /dev/disk0
/dev/disk0: block special
>./a.out /dev/io8log
/dev/io8log: character special
```

# EXERCISE

- Write a C program to simulate echo command
  - Hint: use `read/write` file IO functions.